Name: Manav Jawrani

Roll No.: 19

Subject: Advanced DevOps

Experiment No.: 4

Experiment 4.

Aim: To install kubectl and execute kubectl commands to manage the kubernetes cluster and deploy your First kubernetes Application.

Theory:

- What is kubectl?
kubectl is a kubernetes command-line tool which allows us to run commands against kubernetes clusters. We can use kubectl to deploy applications, inspect and manage cluster resources and view logs.

- Features of kubectl:
1. Communication between Nodes and control plane i.e. master node.
2. User name spaces
3. It allows to customize the outputs
4. It may invoke several API calls to API server to build a view for user

Here we had created our application using nginx web server

- What is nginx?
NGINX is open source software for web serving,

reverse proxying caching load balancing, media streaming and more. It started out as a web server designed for maximum performance and stability.

Deployed our application using Kubernetes deployment

- What is Kubernetes deployment?
A Kubernetes deployment tells Kubernetes how to create or modify instances of the pods that hold a containerized application. Deployments can help to efficiently scale the number of replica pods enable the rollout of updated code in a controlled manner, or roll back to an earlier deployment version if necessary. Kubernetes deployments are completed using Kubectl command-line tool.

Implementation:

****Note: As we have created master and worker nodes and created a kubernetes cluster as well in our previous experiment 3 so will use that only.****

Running An Application on the Cluster

Step 1: As the cluster is up and running, we can deploy our nginx server on this cluster. Apply this deployment file using this command to create a deployment.

$kubectl create deployment nginx --image=nginx

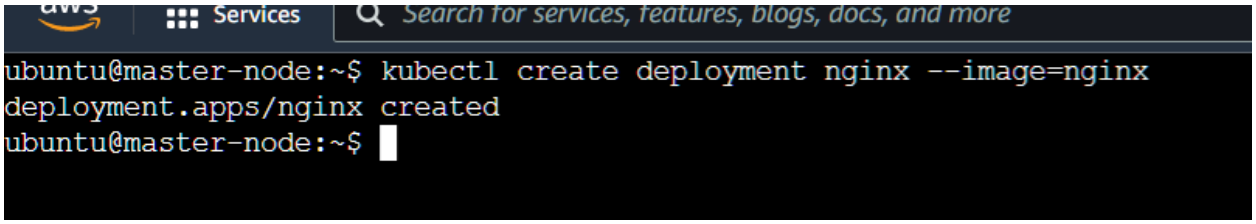

Step 2: Verify the deployment using the command:

$kubectl get deployments

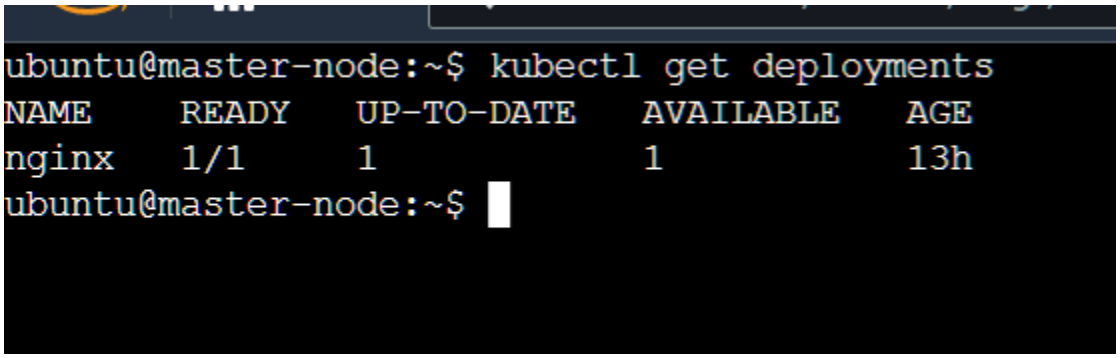

Step 3: Next, run the following command to create a service named nginx that will expose the app publicly. It will do so through a NodePort, a scheme that will make the pod accessible through an arbitrary port opened on each node of the cluster

with this service-type, Kubernetes will assign this service on ports on the **30000+** range.

$kubectl expose deploy nginx --port 80 --target-port 80 --type NodePort

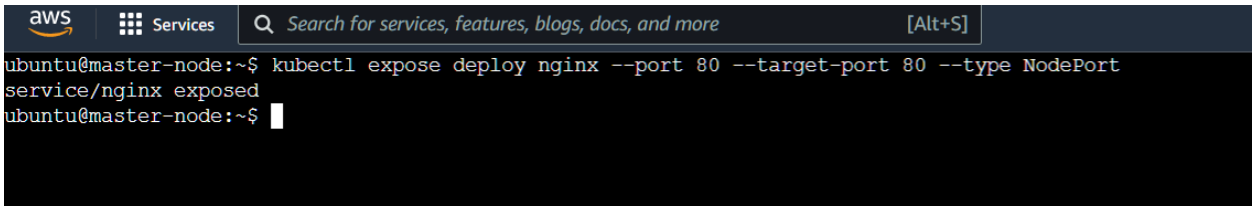

Step 4: Run this command to see a summary of the service and the ports exposed.

$kubectl get services

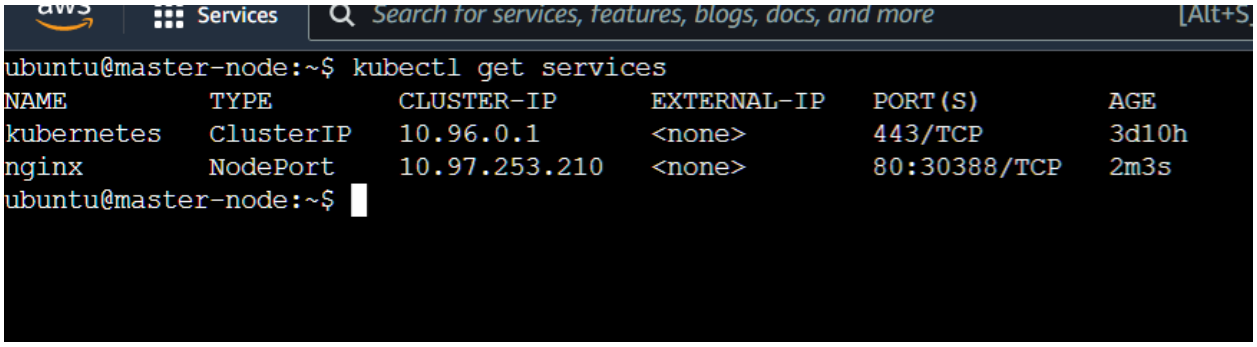

Step 5: Add the port which is displayed i.e. 30388 (in our case) in the inbound rules of the security group.

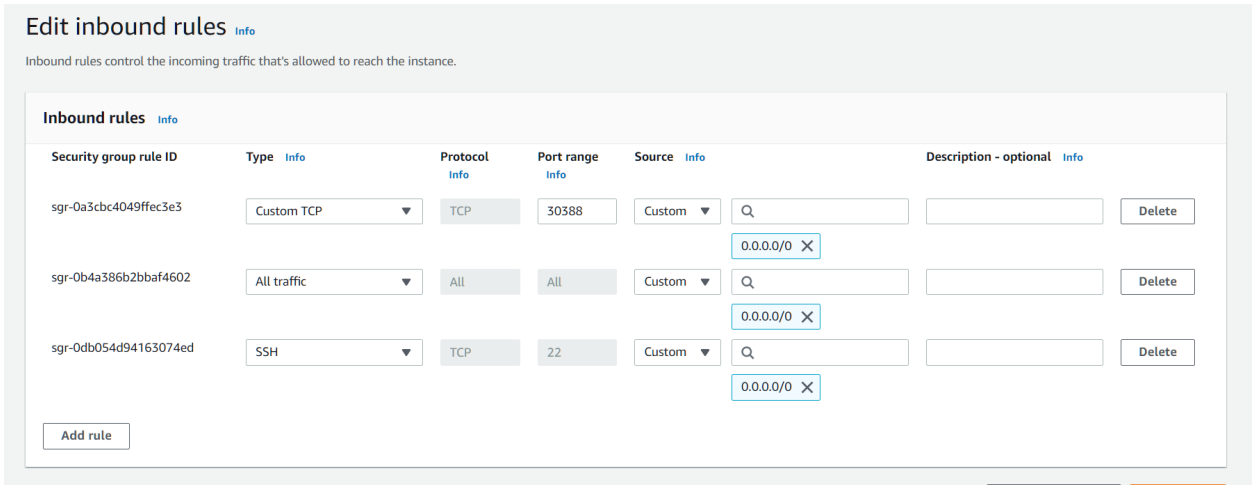

Step 6: Now you can verify that the Nginx page is reachable on all nodes using the curl command.

Master

```
ubuntu@master-node:~$ sudo -i
root@master-node:~# curl master-node:30388
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@master-node:~#
```

Worker 1

```
Reading package lists... Done
root@worker1:~# curl worker1:30388
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@worker1:~#
```

Worker 2

```
root@worker2:~# curl worker2:30388
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
root@worker2:~#
```

As you can see, the "**WELCOME TO NGINX!**" page can be reached.

Step 7: To test that everything is working, visit
http://worker_1_ip:nginx_port
or
http://worker_2_ip:nginx_port
through a browser on your local machine. You will see Nginx's familiar welcome page.

Worker 1: http://52.90.129.234:30388

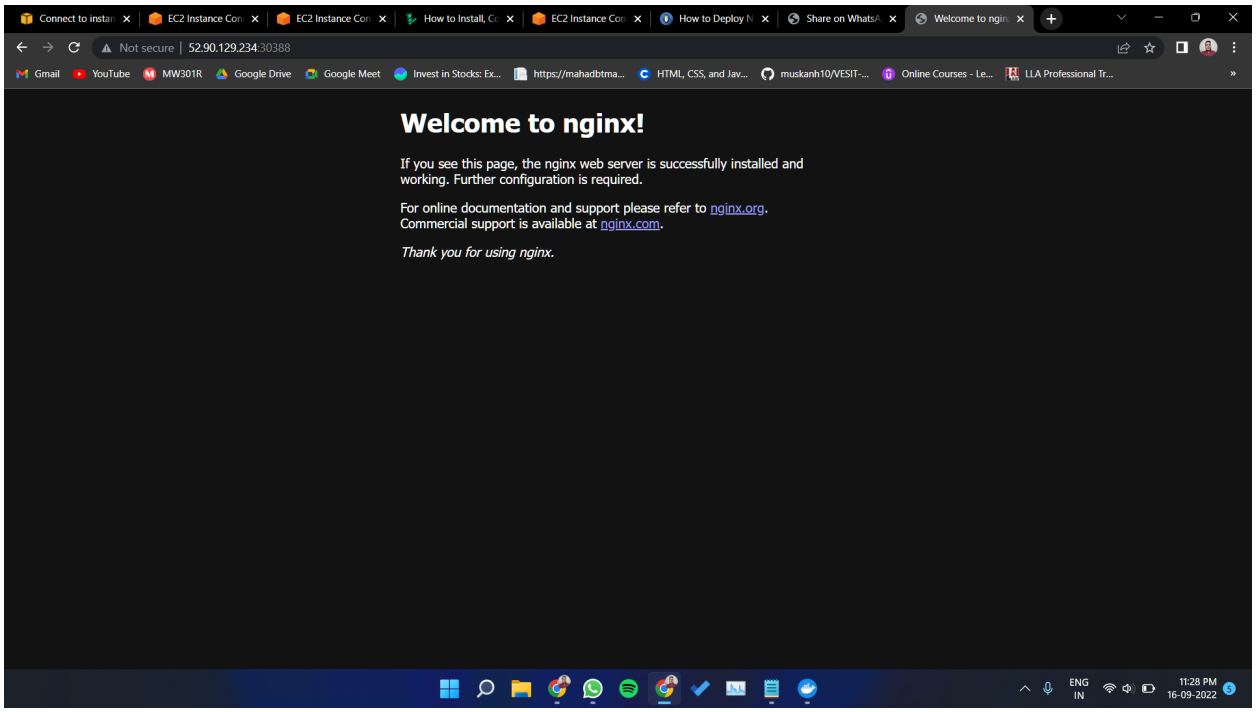

Worker 2: http://3.93.45.75:30388

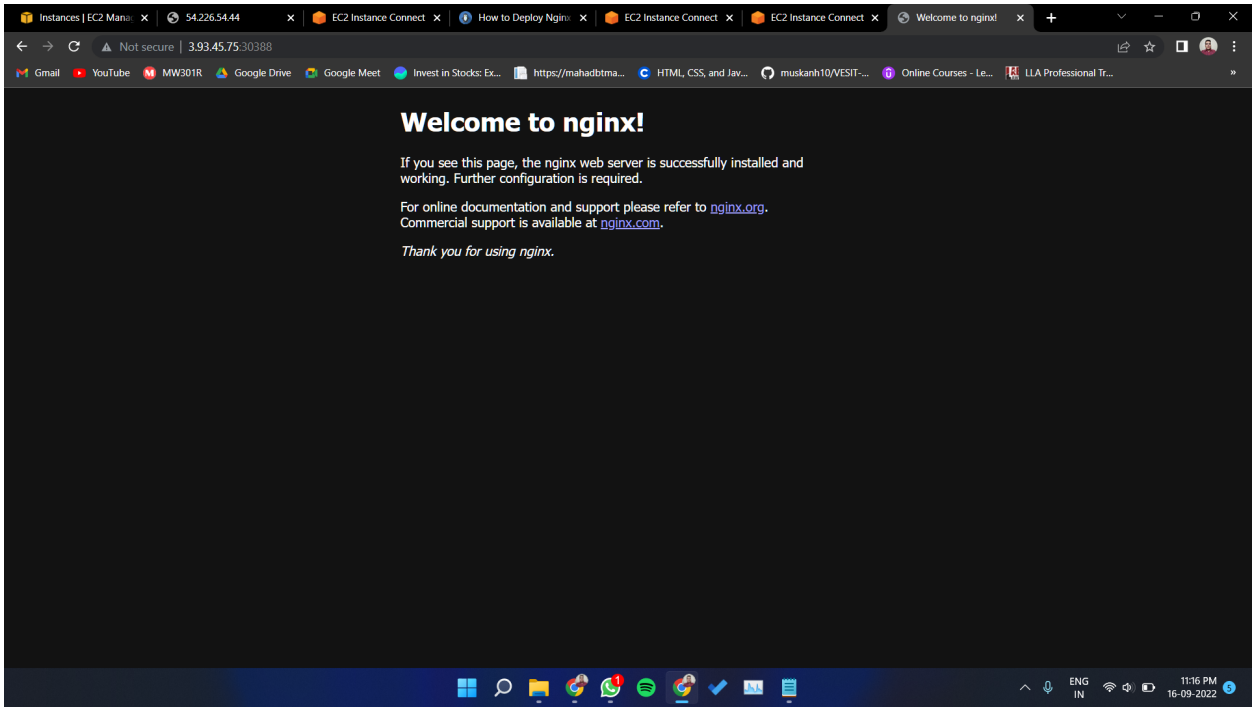

****Note: IPs of the machines will change again and again as soon as you stop the instance and start it again.****

Conclusion :

In this experiment, we have learned how to deploy an application on a Kubernetes cluster using the kubectl command line and Kubernetes deployment, and how to access the application deployment via the cluster's worker nodes.